

# SDSU CS 549 Spring 2024 Machine Learning Lecture 6: PCA

---

IRFAN KHAN

# New Student Assistant/Grader

---

**Student Assistant/Grader:** Manas Gandhi

Email: [mgandhi4512@sdsu.edu](mailto:mgandhi4512@sdsu.edu)

Office location: GMCS 549

Office Hours: TBA

# References

---

SDSU CS549 Lecture Notes by Prof Yang Xu, Spring 2023. Some updated slides used here

Coursera machine learning course by Dr Andrew Ng, Oct 2023

# Outline

---

What is principal component analysis (PCA)? Why do we need it?

Intuitive picture of PCs

Algebraic definition and derivation of PCs


Applications of PCA

# What is Principal Component Analysis (PCA)?

---

- Principal Component Analysis (PCA) is an unsupervised machine learning technique.
- In PCA, the goal is to transform the original features of a dataset into a new set of uncorrelated features called principal components.
- The principal components are ordered by the amount of variance they capture in the data. PCA is commonly used for dimensionality reduction, **visualization**, and noise reduction in datasets.
- Unlike supervised learning, PCA doesn't rely on labeled output for training.
- It analyzes the inherent structure of the data to find patterns and reduce the dimensionality.
- However, PCA can be combined with supervised learning techniques in some applications, such as feature extraction or preprocessing before applying a supervised algorithm

# Example (Car Size and Weight)



	Weight (1000s of Lbs)	Length (feet)	Height (feet)	.....
Car 1	2.0	12	5	
Car 2	6.0	16	6	
Car 3	4.0	14	4.5	
Car 4	5.0	15	5.5	
Car 5	4.5	14	5	
Car 6	3.5	13	4.5	
....				

- Questions:
  - How to visualize?
  - Which features are correlated?
  - Which features are most significant to describe the data?

# Which features are most significant to describe the data?

---

➤ Number of Features  $\Leftrightarrow$  Columns of the Sample Data Matrix ( $m \times n$ )

➤  $m$  is the number of samples

“**Significant**” feature/“dimension of data” should:

➤ Have higher resolution/variability  $\rightarrow$  data points *spread out* across the dimension as much as possible

➤ Be independent  $\rightarrow$  avoid redundancy

# PCA Machine Learning

---

- PCA transforms sample data into a new coordinate system in which:
  - The dimensions are orthogonal (guarantee *independence*) and are
  - ranked according to the variance of data along them (so that more *informative* dimensions, along which the data spread out more, occur first)



# Outline

---


What is principal component analysis (PCA)? Why do we need it?

*Intuitive picture of PCs*

Algebraic definition and derivation of PCs

Applications of PCA


# Sample of $m$ data points in 2-D space



- 2-D Space => 2 features
- **Goal:** to account for the variation in data points with as few dimensions as possible.
- If we are to use only one dimension to describe the data, which one do we choose?
  - $x_1$  or  $x_2$ ?
  - Can we do better - How about other dimensions?

# $m$ Data Points in a 2D-Space

---



$u_1$  and  $u_2$  seem to better represent the data than  $x_1$  and  $x_2$ , as they describe the main “**directions**” of data

# Continuing with Our Car Example Using Only Two Features


	Length (feet) x1	Height (feet) x2
Car 1	12	5
Car 2	16	6
Car 3	14	4.5
Car 4	15	5.5
Car 5	14	5
Car 6	13	4.5
....		

$$\text{Mean } x_1: \mu_1 = \frac{1}{m} \sum_{i=1}^m x_{i1} = 14.0$$

$$\text{Std Dev } x_1: \sigma_1 = \frac{1}{m-1} \sum_{i=1}^m (x_{i1} - \mu_1)^2 = 1.41$$

$$\text{Mean } x_2: \mu_2 = 5.08$$


$$\text{Std Dev } x_2: \sigma_2 = 0.58$$



# Normalizing the Data

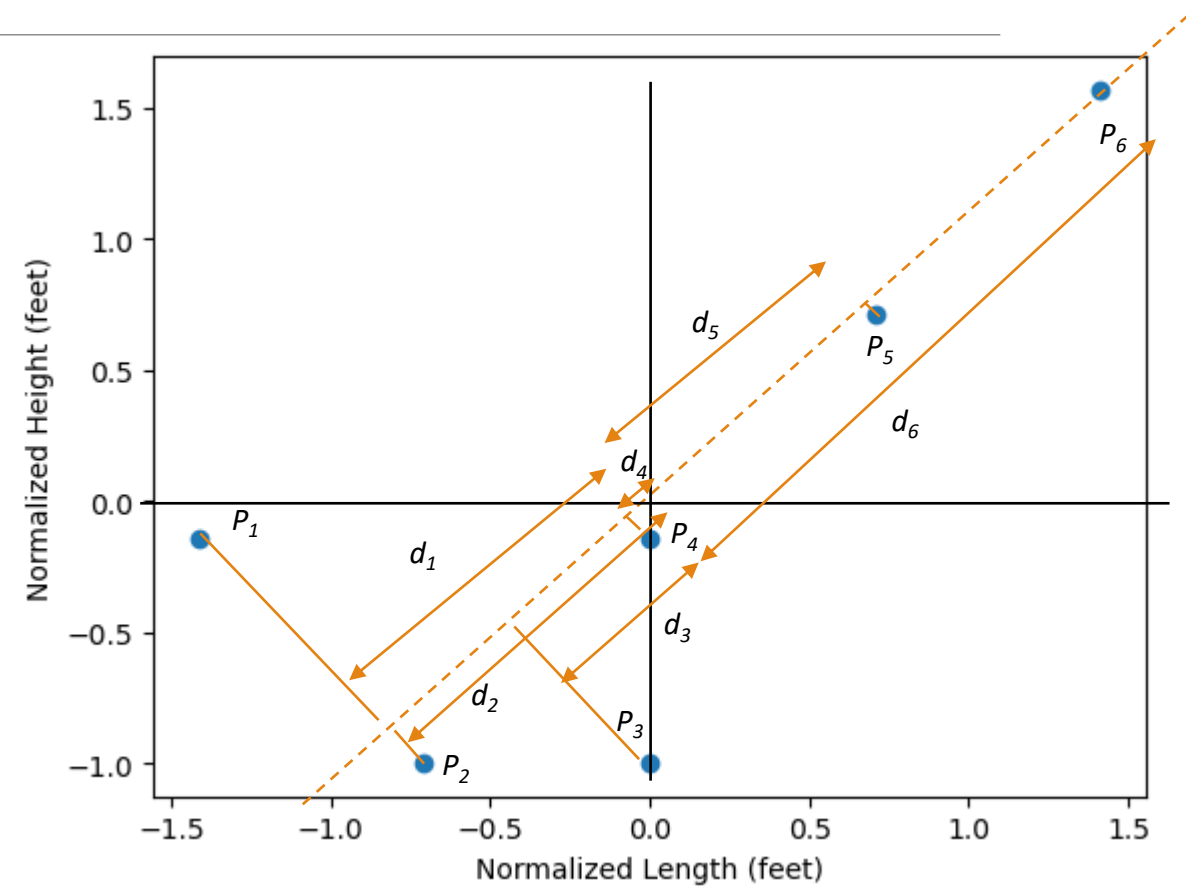
$$\widetilde{x}_1 = \frac{x_1 - \mu_1}{\sigma_1} \quad \widetilde{x}_2 = \frac{x_2 - \mu_2}{\sigma_2}$$

	Normalized Length (feet) ( $\widetilde{x}_1$ )	Normalized Height (feet) ( $\widetilde{x}_2$ )
Car 1	-1.41	-0.14
Car 2	1.41	1.57
Car 3	0	-1.0
Car 4	0.71	0.71
Car 5	0	-0.14
Car 6	-0.71	-1.0




# Quantifying Variance Along an Axis (Dimension)

- Consider the new axis drawn on the scatter plot
- For point  $P_1$ ,  $d_1$  is the projection on the new axis
- Similarly for other points
- Total variance along this dimension represented by the sum of squares of all the projections:
- $SS = \sum_{i=1}^m d_i^2$
- Our objective function to maximize is SS




# Maximize the Sum of Square of Projections



- Which Axis should we choose?
  - One that maximizes  $SS$ :  $PC1$
- Let a unit vector,  $\vec{u}_1$ , represent  $PC1$
- In our example this turns out to be  $\begin{bmatrix} -0.71 \\ 0.71 \end{bmatrix}$
- For each normalized sample data point, the transformed data point, say  $z_1$ , along  $PC1$  is given by the dot product of  $\vec{u}_1$  with  $\tilde{\mathbf{x}} = \begin{bmatrix} \tilde{x}_1 \\ \tilde{x}_2 \end{bmatrix}$
- $z_1 = -0.71 * \tilde{x}_1 + 0.71 * \tilde{x}_2$

# PC1 and PC2


- The second principal component (PC2) is **perpendicular** to PC1
- Let a unit vector,  $\vec{u}_2$ , represent PC2
- In our example this turns out to be  $\begin{bmatrix} -0.71 \\ -0.71 \end{bmatrix}$
- Note that since  $\vec{u}_1$  and  $\vec{u}_2$  are perpendicular, their dot product must be 0.
- For each normalized sample data point, the transformed data point, say  $z_2$ , along PC2 is given by the dot product of  $\vec{u}_2$  with  $\tilde{x} = \begin{bmatrix} \tilde{x}_1 \\ \tilde{x}_2 \end{bmatrix}$ 
  - $z_2 = -0.71 * \tilde{x}_1 - 0.71 * \tilde{x}_2$
- Projections on PC1 have a larger variance than on PC2
  - So  $SS_1 > SS_2$






# Rotate onto the new Axes (basis)

Original Normalized Data



PC1 demonstrates more  
Variance than Length

Represent data using principal components



# Outline

---

What is principal component analysis (PCA)? Why do we need it?

Intuitive picture of PCs

*Algebraic definition and derivation of PCs*

Applications of PCA

# How to compute principal components?

Represent the **principal component (PC)** via a unit length vector  $\vec{u} = \begin{pmatrix} u_1 \\ u_2 \end{pmatrix}$

The length of projection of  $\tilde{x}^{(i)} = \begin{pmatrix} \tilde{x}_1^{(i)} \\ \tilde{x}_2^{(i)} \end{pmatrix}$  along  $\vec{u}$  is the *inner product*:  $\vec{u}^T \tilde{x}^{(i)} = u_1 \tilde{x}_1^{(i)} + u_2 \tilde{x}_2^{(i)}$

Then the sum of square of projections is:

$$SS = \sum_i (\vec{u}^T \tilde{x}^{(i)})^2 = \sum_i (u_1 \tilde{x}_1^{(i)} + u_2 \tilde{x}_2^{(i)})^2 = \sum_i \underbrace{u_1^2 \tilde{x}_1^{(i)2} + 2u_1 u_2 \tilde{x}_1^{(i)} \tilde{x}_2^{(i)} + u_2^2 \tilde{x}_2^{(i)2}}_{\text{Equivalent}}$$


Equivalent

Arrange normalized sample data in a matrix:  $\tilde{X} = \begin{bmatrix} \tilde{x}_1^{(1)} & \tilde{x}_2^{(1)} \\ \vdots & \vdots \\ \tilde{x}_1^{(m)} & \tilde{x}_2^{(m)} \end{bmatrix}$

$m$  is the number of samples

and in general we will have  $n$  features

here  $n$  is 2. So the matrix is  $m \times n$



$$SS = \vec{u}^T (\tilde{X}^T \tilde{X}) \vec{u}$$

Normalized data

Maximize mean SS  $\Rightarrow$  maximize  $\frac{1}{m} SS = \vec{u}^T \left( \frac{1}{m} \tilde{X}^T \tilde{X} \right) \vec{u}$

Under constraint  $\|\vec{u}\| = 1$

# How to compute principal components (cont.)

A typical constrained optimization problem:

$$\text{Maximize } \frac{1}{m} SS = \vec{u}^T \left( \frac{1}{m} \tilde{X}^T \tilde{X} \right) \vec{u}, \text{ with } \|\vec{u}\| = 1$$

Measures how much  $x_1$  and  $x_2$  are related (move in the same direction)



$$\frac{1}{m} \tilde{X}^T \tilde{X} = \frac{1}{m} \begin{bmatrix} \frac{x_1^{(1)} - \mu_1}{\sigma_1} & \dots & \frac{x_1^{(m)} - \mu_1}{\sigma_1} \\ \frac{x_2^{(1)} - \mu_2}{\sigma_2} & \dots & \frac{x_2^{(m)} - \mu_2}{\sigma_2} \end{bmatrix} \begin{bmatrix} \frac{x_1^{(1)} - \mu_1}{\sigma_1} & \frac{x_2^{(1)} - \mu_2}{\sigma_2} \\ \dots & \dots \\ \frac{x_1^{(m)} - \mu_1}{\sigma_1} & \frac{x_2^{(m)} - \mu_2}{\sigma_2} \end{bmatrix} = \frac{1}{n} \begin{bmatrix} \boxed{\sum \frac{(x_1^{(i)} - \mu_1)^2}{\sigma_1^2}} & \boxed{\sum \frac{(x_1^{(i)} - \mu_1)(x_2^{(i)} - \mu_2)}{\sigma_1 \sigma_2}} \\ \boxed{\sum \frac{(x_1^{(i)} - \mu_1)(x_2^{(i)} - \mu_2)}{\sigma_1 \sigma_2}} & \boxed{\sum \frac{(x_2^{(i)} - \mu_2)^2}{\sigma_2^2}} \end{bmatrix}$$

variance of  $x_1$       covariance between  $x_1$  and  $x_2$

variance of  $x_2$

To solve this maximization problem, use the **Lagrange Multiplier** method:

$2 \times 2$  covariance matrix for  $n=2$  (in general  $n \times n$ )

Sometimes denoted as  $\Sigma$

$$\text{maximize (unconstrained): } L = \vec{u}^T \left( \frac{1}{m} \tilde{X}^T \tilde{X} \right) \vec{u} - \lambda (\vec{u}^T \vec{u} - 1)$$

$$\frac{\partial L}{\partial \vec{u}} = \left( \frac{1}{m} \tilde{X}^T \tilde{X} \right) \vec{u} - \lambda \vec{u} = 0 \implies \left( \frac{1}{m} \tilde{X}^T \tilde{X} \right) \vec{u} = \lambda \vec{u}$$

Solution  $\vec{u}$  is the principal **eigenvector** of  $\frac{1}{m} \tilde{X}^T \tilde{X}$

The maximized SS is the corresponding eigenvalue, representing variance along the Eigen-Vector

# Linear Algebra Review

“Eigen” is the Dutch word meaning “my”

## **Eigenvectors and eigenvalues**

---

Eigenvectors  $\approx$  “my vectors”

For square matrix  $A$  and column vector  $\vec{v}$ , if

$$A\vec{v} = \lambda\vec{v}$$

Then  $\vec{v}$  is the eigenvector (or characteristic vector) of  $A$ , and  $\lambda$  is the eigenvalue (or characteristic value)

### **Meanings:**

$A$  is a linear transformation. Almost all vectors change direction when they are transformed by  $A$  (i.e., multiplied by  $A$ )

Certain exceptional vectors  $\vec{v}$  are in the same direction of  $A\vec{v}$ . Those are eigenvectors.  
Transform an eigenvector by  $A$ , the resulting vector  $A\vec{v}$  is a number  $\lambda$  times the original  $\vec{v}$ .

# Linear algebra reviews (cont.)

## **Spectral theorem:**

If the  $n \times n$  matrix  $A$  is *symmetric* ( $A^T = A$ ), then  $A$  is orthogonally diagonalizable and has only real eigenvalues.

In other words, there exist real numbers  $\lambda_1, \dots, \lambda_n$  (**eigenvalues**) and orthogonal, non-zero real vectors  $\vec{v}_1, \dots, \vec{v}_n$  (**eigenvectors**) such that for each  $i = 1, 2, \dots, n$ :

$$A\vec{v}_i = \lambda_i \vec{v}_i$$

If  $\tilde{X}$  is a  $m \times n$  matrix, then the  $n \times n$  matrix  $\tilde{X}^T \tilde{X}$  and the  $m \times m$  matrix  $\tilde{X} \tilde{X}^T$  are both symmetric.

The matrices  $\frac{1}{m} \tilde{X} \tilde{X}^T$  and  $\frac{1}{m} \tilde{X}^T \tilde{X}$  share the same nonzero eigenvalues.

*Proof:* let  $\vec{v}$  be a nonzero eigenvector of  $\frac{1}{m} \tilde{X} \tilde{X}^T$  with eigenvalue  $\lambda \neq 0$

$$\left( \frac{1}{m} \tilde{X} \tilde{X}^T \right) \vec{v} = \lambda \vec{v}$$

Multiply both sides with  $\tilde{X}^T$  on the left:

$$\tilde{X}^T \left( \frac{1}{m} \tilde{X} \tilde{X}^T \right) \vec{u} = \tilde{X}^T \lambda \vec{v} \implies \underbrace{\left( \frac{1}{m} \tilde{X}^T \tilde{X} \right)}_{\text{eigenvector}} (\tilde{X}^T \vec{v}) = \lambda (\tilde{X}^T \vec{v})$$

Usage:

if  $m$  and  $n$  are very different, say,  $X$  is  $2 \times 500$

Then the quick way to find the eigenvalues of the  $500 \times 500$  matrix  $X^T X$  is to first find the 2 eigenvalues of the  $2 \times 2$  matrix  $XX^T$ .

The other 498 eigenvalues of  $X^T X$  are all zero!

# Eigenvalues and eigenvectors of covariance matrix

When  $X$  is  $m \times 2$

$$\Sigma = \frac{1}{n} \tilde{X}^T \tilde{X} = \begin{bmatrix} \text{Var}(x_1) & \text{Cov}(x_1, x_2) \\ \text{Cov}(x_1, x_2) & \text{Var}(x_2) \end{bmatrix} \text{ is a symmetric } 2 \times 2 \text{ matrix} \Rightarrow \begin{matrix} \text{There are 2 eigenvectors and eigenvalues} \\ \Sigma u_1 = \lambda_1 u_1 & \Sigma u_2 = \lambda_2 u_2 & \text{Let } \lambda_1 > \lambda_2 \end{matrix}$$

Total variance in data  $T = \lambda_1 + \lambda_2$       PC1  $u_1$  accounts for  $\frac{\lambda_1}{T}$  of the total variance      PC2  $u_2$  accounts for  $\frac{\lambda_2}{T}$

---

What if  $X$  is  $m \times n \rightarrow$  high dimensional data with  $n$  features

$$\Sigma = \frac{1}{n} \tilde{X}^T \tilde{X} = \begin{bmatrix} \text{Var}(x_1) & \text{Cov}(x_1, x_2) & \dots & \text{Cov}(x_1, x_n) \\ \text{Cov}(x_1, x_2) & \text{Var}(x_2) & \dots & \text{Cov}(x_2, x_n) \\ \vdots & \vdots & \ddots & \vdots \\ \text{Cov}(x_1, x_n) & \text{Cov}(x_2, x_n) & \dots & \text{Var}(x_n) \end{bmatrix} \text{ is a symmetric } n \times n \text{ matrix} \Rightarrow \begin{matrix} \text{There are } n \text{ eigenvectors and eigenvalues} \\ \Sigma u_1 = \lambda_1 u_1 & \Sigma u_2 = \lambda_2 u_2 & \dots & \Sigma u_n = \lambda_n u_n \\ \text{Let } \lambda_1 > \lambda_2 > \dots > \lambda_n \end{matrix}$$

Total variance in data  $T = \lambda_1 + \lambda_2 + \dots + \lambda_n$       PC1  $u_1$  accounts for  $\frac{\lambda_1}{T}$  of the total variance      PC2  $u_2$  accounts for  $\frac{\lambda_2}{T}$

...      PCn  $u_n$  accounts for  $\frac{\lambda_n}{T}$

# Finding *eigenvalues* and *eigenvectors* in Python

---

Use the `linalg` module of `numpy`

```
► In [2]: from numpy import linalg as LA  
import numpy as np
```

```
In [3]: X = np.array([[10,11,8,3,2,1], [6,4,5,3,2.8,1]])
```

```
In [6]: w, v = LA.eig(np.dot(X, X.T))
```

```
In [7]: print(w)  
print(v)
```

```
[386.37312455  7.46687545]  
[[ 0.87715849 -0.48020099]  
 [ 0.48020099  0.87715849]]
```



# What's next? Reduce to $k$ dimensions

---

$X$  is  $m \times n$       $\Sigma = \frac{1}{m} \tilde{X}^T \tilde{X}$  is  $m \times m$      eigenvectors and eigenvalues:  $\lambda_i, u_i$

Select the top  $k$  eigenvectors  $u_1, u_2, \dots, u_k$ , with top  $k$  eigenvalues  $\lambda_1 > \lambda_2 > \dots > \lambda_k$

$$\begin{array}{ccc}
 \mathbf{U} = [\mathbf{u}_1 \ \mathbf{u}_2 \ \dots \ \mathbf{u}_k] & \text{Span to a } k\text{-dimensional subspace} & \Rightarrow \text{Project data point } \tilde{\mathbf{x}}^{(i)} \text{ into this subspace} \\
 n \times k \text{ matrix} & & \mathbf{y}^{(i)} = \underbrace{\tilde{\mathbf{x}}^{(i)}}_{1 \times n} \mathbf{U} = [y_1^{(i)} \ y_2^{(i)} \ \dots \ y_k^{(i)}] \\
 & & \underbrace{\hspace{1.5cm}}_{1 \times k}
 \end{array}$$

Project  $m$  examples into new subspace

$$\underbrace{\mathbf{Y}}_{m \times k} = \underbrace{\tilde{\mathbf{X}}}_{m \times n} \mathbf{U} \quad n \times k$$

Reduce  $n$ -dimensional data to  $k$ -dimensional

$\mathbf{Y}$  are the coordinates of data in the new subspace

For visualization,  $k = 2$  or  $3$

# Reconstruction of Original Data from PCs

---

- Let  $\tilde{\mathbf{x}}$  be the original data and its projections on the first  $k$  PCs ( $\mathbf{u}_1, \mathbf{u}_2, \dots, \mathbf{u}_k$ ) are given by:
  - $\mathbf{y}_j = \tilde{\mathbf{x}} \mathbf{u}_j$  for  $j = 1, 2, \dots, k$
- Reconstructed data  $\tilde{\mathbf{x}}'$  is given by:
  - $\tilde{\mathbf{x}}' = \sum_{j=1}^k \mathbf{y}_j \mathbf{u}_j^\top$
- If the full set of PCs are used, then the reconstruction will be perfect, i.e., exactly the same as the original image without losing any information.
- If a subset (e.g., top  $k$  PCs) is used, then the reconstruction will cause some information loss.
- This information loss can be measured by the Euclidean distance between the original data  $\tilde{\mathbf{x}}$  and the reconstructed data  $\tilde{\mathbf{x}}'$ . Larger distance indicates higher information loss.

# Covariance matrix of the transformed data

$Y$  is the transformed data



Its covariance matrix:  $\Sigma = \frac{1}{m} Y^T Y = \frac{1}{m} (\tilde{X} U)^T (\tilde{X} U) = \frac{1}{m} U^T \tilde{X}^T \tilde{X} U = \begin{bmatrix} \lambda_1 & & \\ & \lambda_2 & \\ & & \ddots \\ & & & \lambda_k \end{bmatrix}$

A diagonal matrix with eigenvalues as diagonal elements

Therefore, all covariances  $Var(y_i, y_j)$  are **zero!** → All dimensions are independent

---

Covariance matrix and the shape of data



# Summary: procedures of performing PCA

---

Normalize data:  $X \rightarrow \tilde{X}, m \times n$

Find the eigenvectors and eigenvalues of the covariance matrix  $\frac{1}{m} \tilde{X}^T \tilde{X}, u_1, u_2, \dots, u_n, \lambda_1 > \lambda_2 > \dots > \lambda_n$

Observe whether a small number of the  $\lambda_i$  are much bigger than all the others, e.g., first  $k$  elements.

- If yes, a dimension reduction is proper:  $n\text{-dim} \rightarrow k\text{-dim}$

Interpret the principal components:

- Which ones are the most important?

# Applications of PCA

- *Better interpretation of features*

Case 1: Sibley's Bird Database of North American birds

	Length, $x_1$	Wingspan, $x_2$	Weight, $x_3$
Bird 1	4	10	6
Bird 2	5	11	4
Bird 3	6	8	5
...	...	...	...

$$\lambda_1 = 1626.5 \quad \lambda_2 = 129.0 \quad \lambda_3 = 7.1$$

$$u_1 = \begin{bmatrix} 0.22 \\ 0.41 \\ 0.88 \end{bmatrix} \quad u_2 = \begin{bmatrix} 0.25 \\ 0.85 \\ -0.46 \end{bmatrix} \quad u_3 = \begin{bmatrix} 0.94 \\ -0.32 \\ -0.08 \end{bmatrix}$$

$X$  is  $100 \times 3$

$$\Sigma = \begin{bmatrix} 91.4 & 171.9 & 298.0 \\ & 373.9 & 545.2 \\ & & 1297.3 \end{bmatrix}$$

$$\frac{\lambda_1}{\lambda_1 + \lambda_2 + \lambda_3} = 92.3\% \quad \frac{\lambda_2}{\lambda_1 + \lambda_2 + \lambda_3} = 7.3\%$$

$$\text{PC1} = 0.22 \text{ Length} + 0.41 \text{ Wingspan} + 0.88 \text{ Weight}$$

It characterizes the “size”

$$\text{PC2} = 0.25 \text{ Length} + 0.85 \text{ Wingspan} - 0.46 \text{ Weight}$$


It characterizes the “stoutness”

# Applications of PCA


## Data compression

- Original data:  $X \in \mathbb{R}^n$
- Use the top  $k$  PCs to approximate:  $X \cong X^k$


Original image with 784 dimensions



Compressed image with 184 dimensions



Compressed image with 10 components



Source: <https://towardsdatascience.com/image-compression-using-principal-component-analysis-pca-253f26740a9f>

# Applications of PCA: Eigenface

---

A training set of  $m$  face images. Each image is  $100 \times 100$

Flatten each image to column vector and stack them  $\rightarrow X$  of shape  $m \times 10000$  (normalized)

Compute the eigenvectors and eigenvalues of the covariance matrix  $\frac{1}{m} X^T X$ , of shape  $10000 \times 10000$

- Each eigenvector is of shape  $10000 \times 1$
- Reshape it back to  $100 \times 100 \rightarrow$  It can be seen as a face image: “eigenface”

These 10000 eigenfaces can be used to represent new faces

- Project a new image vector to a selected subset of eigenfaces
- The projections can be used to identify the new face
- In practice, choosing **100 to 150** eigenfaces are sufficient

Example of the first eigenvector (PC1)  
From Jauregui (2012)



# Assignment

---

- Assignment uses sign-language digits data – provided in a numpy array
- The original data can be obtained from <https://www.kaggle.com/ardamavi/sign-language-digits-dataset> (preprocessed) or <https://github.com/ardamavi/Sign-Language-Digits-Dataset> (raw).